

Exercise 2 and 3

November 26, 2021

Machine Learning Lab

Lab 03

0.1 Exercise 2: Linear Regression with Gradient Descent

0.1.1 Part A: (Datasets)

Importing Packages

```
[1]: import pandas as pd          #Importing Pandas
import numpy as np              #Importing Numpy
import matplotlib.pyplot as plt #Importing Matplotlib
```

Reading Wine Quality Dataset

```
[2]: wine_quality = pd.read_csv('winequality-red.csv',delimiter=';')
wine_quality.head()
```

```
[2]:   fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4             0.70         0.00           1.9        0.076
1           7.8             0.88         0.00           2.6        0.098
2           7.8             0.76         0.04           2.3        0.092
3          11.2             0.28         0.56           1.9        0.075
4           7.4             0.70         0.00           1.9        0.076
```

```
   free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0             11.0             34.0  0.9978  3.51         0.56
1             25.0             67.0  0.9968  3.20         0.68
2             15.0             54.0  0.9970  3.26         0.65
3             17.0             60.0  0.9980  3.16         0.58
4             11.0             34.0  0.9978  3.51         0.56
```

```
   alcohol  quality
0       9.4        5
1       9.8        5
2       9.8        5
3       9.8        6
4       9.4        5
```

Reading Airfare and Demand Dataset

```
[3]: airfare_demand = pd.read_fwf('airq402.data',header=None)
airfare_demand.columns = ['City1','City2','Average Fare1','Distance','Average_
↪weekly passengers',
                           'market leading airline','market share','Average_
↪fare2','Low price airline',
                           'market share','price']
airfare_demand.head()
```

```
[3]:  City1 City2  Average Fare1  Distance  Average weekly passengers \
0    CAK   ATL         114.47      528                424.56
1    CAK   MCO         122.47      860                276.84
2    ALB   ATL         214.42      852                215.76
3    ALB   BWI          69.40      288                606.84
4    ALB   ORD         158.13      723                313.04

    market leading airline  market share  Average fare2 Low price airline \
0                        FL          70.19         111.03              FL
1                        FL          75.10         123.09              DL
2                        DL          78.89         223.98              CO
3                        WN          96.97          68.86              WN
4                        UA          39.79         161.36              WN

    market share  price
0          70.19  111.03
1          17.23  118.94
2           2.77  167.12
3          96.97   68.86
4          15.34  145.42
```

Reading Parkison Dataset

```
[4]: parkison = pd.read_csv('parkinsons_updrs.data')
parkison.head()
```

```
[4]:  subject#  age  sex  test_time  motor_UPDRS  total_UPDRS  Jitter(%) \
0         1   72    0    5.6431      28.199      34.398    0.00662
1         1   72    0   12.6660      28.447      34.894    0.00300
2         1   72    0   19.6810      28.695      35.389    0.00481
3         1   72    0   25.6470      28.905      35.810    0.00528
4         1   72    0   33.6420      29.187      36.375    0.00335

    Jitter(Abs)  Jitter:RAP  Jitter:PPQ5  ...  Shimmer(dB)  Shimmer:APQ3 \
0    0.000034    0.00401    0.00317  ...    0.230      0.01438
1    0.000017    0.00132    0.00150  ...    0.179      0.00994
2    0.000025    0.00205    0.00208  ...    0.181      0.00734
3    0.000027    0.00191    0.00264  ...    0.327      0.01106
4    0.000020    0.00093    0.00130  ...    0.176      0.00679
```

	Shimmer:APQ5	Shimmer:APQ11	Shimmer:DDA	NHR	HNR	RPDE	\
0	0.01309	0.01662	0.04314	0.014290	21.640	0.41888	
1	0.01072	0.01689	0.02982	0.011112	27.183	0.43493	
2	0.00844	0.01458	0.02202	0.020220	23.047	0.46222	
3	0.01265	0.01963	0.03317	0.027837	24.445	0.48730	
4	0.00929	0.01819	0.02036	0.011625	26.126	0.47188	

	DFA	PPE
0	0.54842	0.16006
1	0.56477	0.10810
2	0.54405	0.21014
3	0.57794	0.33277
4	0.56122	0.19361

[5 rows x 22 columns]

Convert any non-numeric values to numeric values. For example you can replace a country name with an integer value or more appropriately use hot-one encoding. [Hint: use `pandas.get_dummies`]. Please explain your solution.

```
[5]: wine_quality.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

We can see from the Wine Quality Dataset info that No Column is Non-Numeric, so we skip Hot One Encoding here

```
[6]: airfare_demand.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
```

Data columns (total 11 columns):

#	Column	Non-Null Count	Dtype
0	City1	1000 non-null	object
1	City2	1000 non-null	object
2	Average Fare1	1000 non-null	float64
3	Distance	1000 non-null	int64
4	Average weekly passengers	1000 non-null	float64
5	market leading airline	1000 non-null	object
6	market share	1000 non-null	float64
7	Average fare2	1000 non-null	float64
8	Low price airline	1000 non-null	object
9	market share	1000 non-null	float64
10	price	1000 non-null	float64

dtypes: float64(6), int64(1), object(4)

memory usage: 86.1+ KB

Here we see that the Columns : 0,1,5,8 are Non-Numeric Columns, so converting these columns into suitable Hot one Encoding

```
[7]: #Converting Non-Numeric Columns to to numeric Columns using Hot one encoding
      ↪scheme
airfare_demand_encoded = pd.
      ↪get_dummies(airfare_demand,columns=['City1','City2','market leading_
      ↪airline','Low price airline'])
airfare_demand_encoded.head()
```

```
[7]:   Average Fare1  Distance  Average weekly passengers  market share \
0          114.47      528                424.56          70.19
1          122.47      860                276.84          75.10
2          214.42      852                215.76          78.89
3           69.40      288                606.84          96.97
4          158.13      723                313.04          39.79

   Average fare2  market share  price  City1_ABQ  City1_ACY  City1_ALB  ... \
0          111.03          70.19  111.03         0         0         0  ...
1          123.09          17.23  118.94         0         0         0  ...
2          223.98           2.77  167.12         0         0         1  ...
3           68.86          96.97   68.86         0         0         1  ...
4          161.36          15.34  145.42         0         0         1  ...

   Low price airline_G4  Low price airline_HP  Low price airline_NK \
0                     0                     0                     0
1                     0                     0                     0
2                     0                     0                     0
3                     0                     0                     0
4                     0                     0                     0
```

	Low price airline_NW	Low price airline_SY	Low price airline_TZ \
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

	Low price airline_UA	Low price airline_US	Low price airline_WN \
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	1
4	0	0	1

	Low price airline_YX
0	0
1	0
2	0
3	0
4	0

[5 rows x 217 columns]

```
[8]: parkison.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5875 entries, 0 to 5874
Data columns (total 22 columns):
#   Column                Non-Null Count  Dtype
---  -
0   subject#              5875 non-null   int64
1   age                   5875 non-null   int64
2   sex                   5875 non-null   int64
3   test_time             5875 non-null   float64
4   motor_UPDRS           5875 non-null   float64
5   total_UPDRS           5875 non-null   float64
6   Jitter(%)            5875 non-null   float64
7   Jitter(Abs)           5875 non-null   float64
8   Jitter:RAP            5875 non-null   float64
9   Jitter:PPQ5           5875 non-null   float64
10  Jitter:DDP            5875 non-null   float64
11  Shimmer               5875 non-null   float64
12  Shimmer(dB)           5875 non-null   float64
13  Shimmer:APQ3          5875 non-null   float64
14  Shimmer:APQ5          5875 non-null   float64
15  Shimmer:APQ11         5875 non-null   float64
16  Shimmer:DDA           5875 non-null   float64
17  NHR                   5875 non-null   float64
```

```

18  HNR          5875 non-null   float64
19  RPDE         5875 non-null   float64
20  DFA          5875 non-null   float64
21  PPE          5875 non-null   float64
dtypes: float64(19), int64(3)
memory usage: 1009.9 KB

```

We can see from the Parkisons Dataset info that No Column is Non-Numeric, so we skip Hot One Encoding here

If required drop out the rows with missing values or NA. In next lectures we will handle sparse data, which will allow us to use records with missing values.

```
[9]: wine_quality.isnull().sum()
```

```

[9]: fixed acidity          0
    volatile acidity        0
    citric acid             0
    residual sugar          0
    chlorides               0
    free sulfur dioxide      0
    total sulfur dioxide     0
    density                 0
    pH                     0
    sulphates               0
    alcohol                 0
    quality                 0
    dtype: int64

```

There are No Empty and NA values in Wine Quality Dataset

```
[10]: airfare_demand_encoded.isnull().sum()
```

```

[10]: Average Fare1          0
    Distance                 0
    Average weekly passengers 0
    market share             0
    Average fare2            0
    ..
    Low price airline_TZ      0
    Low price airline_UA      0
    Low price airline_US      0
    Low price airline_WN      0
    Low price airline_YX      0
    Length: 217, dtype: int64

```

There are also No Empty and NA values in Airfare Demand Dataset

```
[11]: parkison.isnull().sum()
```

```
[11]: subject#      0
      age          0
      sex          0
      test_time    0
      motor_UPDRS  0
      total_UPDRS  0
      Jitter(%)    0
      Jitter(Abs)  0
      Jitter:RAP    0
      Jitter:PPQ5   0
      Jitter:DDP    0
      Shimmer      0
      Shimmer(dB)   0
      Shimmer:APQ3  0
      Shimmer:APQ5  0
      Shimmer:APQ11 0
      Shimmer:DDA   0
      NHR          0
      HNR          0
      RPDE         0
      DFA          0
      PPE          0
      dtype: int64
```

There are also No Empty and NA values in Parkisons Dataset

Split the dataset into 80% Train set and 20% Test set.

Function to Split any Dataframe into Training Set and Test Set

```
[12]: def split_dataset(dataset,label):
      #Creating X matrix by removing the label/Target column
      X = dataset.drop(label,axis=1)
      #Creating Y vector which include only the Label/Target Column
      Y = dataset[label].to_numpy()

      #Adding a Bias Column in the X Matrix
      X = np.append(np.ones(shape=(len(X),1)),X,axis=1)

      #Calculating number of rows to be copied in the Training Set according to
      ↪80:20 ratio
      total_training_rows = int(len(X)*0.8)

      #Splitting the Dataset into Training set and Test Set based on calculated
      ↪rows
      X_train , Y_train = X[:total_training_rows,:] , Y[:total_training_rows].
      ↪reshape(-1,1)
```

```

    X_test , Y_test = X[total_training_rows:,:] , Y[total_training_rows:].
    ↪reshape(-1,1)

    return X_train,X_test,Y_train,Y_test

```

Splitting Wine Quality Dataset

```

[13]: wine_Xtrain , wine_Xtest , wine_Ytrain , wine_Ytest = ↪
    ↪split_dataset(wine_quality,'quality')

```

Splitting Airfare Demand Dataset

```

[14]: airfare_Xtrain , airfare_Xtest , airfare_Ytrain , airfare_Ytest = ↪
    ↪split_dataset(airfare_demand_encoded,'price')

```

Splitting Parkinson Dataset

```

[15]: parkison_Xtrain , parkison_Xtest , parkison_Ytrain , parkison_Ytest = ↪
    ↪split_dataset(parkison,'total_UPDRS')

```

0.1.2 Part B: Linear Regression with Real-World Data

```

[16]: #Initializing arrays to store loss difference and RMSE values in different ↪
    ↪number of Iterations
    loss_difference_values = np.array([])
    rmse_values = np.array([])

```

Function to Calculate Loss Between Actual Y and Predicted Y

```

[17]: def loss_function(X,Y,B):
    # L(B) = summation((y-ypred)^2)
    return np.sum(np.square(np.subtract(Y,X @ B)))

```

Function to Calculate the Loss Difference based on Old Beta Values and New Beta Values

```

[18]: def loss_difference(X,Y,B_old,B_new):
    # |L(B_old) - L(B_new)|
    return np.abs(loss_function(X,Y,B_old) - loss_function(X,Y,B_new))

```

Function to Calculate the RMSE Loss between Actual Y and Predicted Y

```

[19]: def rmse(X,Y,B):
    # RMSE = square_root(summation((y-ypred)^2)/N)
    return np.sqrt(np.sum(np.square(np.subtract(Y , X @ B)))/len(X))

```

Function which returns the gradient of Loss function

```

[20]: def dL(X,Y,B):
    # Derivation of Loss = -2 * X^t * (y - ypred)
    return -2 * (X.T @ (Y - X @ B))

```

Function to Minimize Gradient Descent based on Total Iterations and Learning Rate


```
[21]: def minimize_GD(X,Y,X_test,Y_test,imax,mu):
    #Using global arrays to store loss difference and rmse values for different
    Iterations
    global loss_difference_values , rmse_values

    #Emptying both Loss difference and RMSE arrays
    loss_difference_values , rmse_values = np.array([]) , np.array([])

    #Initializing beta with Zeros
    beta = np.zeros(shape=(len(X[0]),1))
    for i in range(imax):
        #Calculating new Beta values from previous beta values and gradient
        descent direction
        #Beta = Beta - learning_rate * gradient Descent based on Beta
        beta_ = beta - mu * dL(X,Y,beta)

        #Appending Loss difference between between previous and new Beta
        loss_difference_values = np.
        append(loss_difference_values,loss_difference(X,Y,beta,beta_))

        #Appending RMSE loss between actual Y and Predicted Y
        rmse_values = np.append(rmse_values,rmse(X_test,Y_test,beta_))

        #Copying new Beta value to old Beta value for Further Calculation
        beta = np.copy(beta_)
    return beta_
```

Function to Plot Loss Difference and RMSE Loss

```
[22]: def plot_loss(xvalues , loss_difference , rmse ,title_graph1 ,title_graph2,
    xlabel_graph1 ,xlabel_graph2 ,ylabel_graph1,ylabel_graph2):
    #Plotting Loss Difference Graph
    fig = plt.figure(figsize=(15,5))
    ax = fig.add_subplot(121)
    ax.plot(xvalues,loss_difference)
    ax.set_title(title_graph1)
    ax.set_xlabel(xlabel_graph1)
    ax.set_ylabel(ylabel_graph1)

    #Plotting RMSE Loss
    ax1 = fig.add_subplot(122)
    ax1.plot(xvalues,rmse)
    ax1.set_title(title_graph2)
    ax1.set_xlabel(xlabel_graph2)
    ax1.set_ylabel(ylabel_graph2)
```

Applying Different Iteration and Different mu values for the Optimum values


```
[23]: imax_d = [100,500,1000]
mu_wine = [0.00000001,0.00000001,0.000000001]
mu_parkison = [0.00000001,0.000000001,0.0000000001]
mu_airfare = [0.000000000001,0.0000000000001,0.00000000000001]
```

Applying Different Iterations on Wine Quality Dataset with $\mu=0.0000001$


```
[24]: for i in imax_d:
    #Calculating new beta using Gradient Descent using  $\mu = 0.0000001$ 
    beta = minimize_GD(wine_Xtrain,wine_Ytrain,wine_Xtest,wine_Ytest,i,0.
    ↪0000001)

    #Plotting the Loss and RMSE with different Iterations
    plot_loss([j for j in range(i)],loss_difference_values,rmse_values,
    ↪'Wine Quality Dataset | Graph of  $|f(x_{i-1})-f(x_i)|$  with Respect to_
    ↪Iterations\n'.format(i),
    ↪'Wine Quality Dataset | Graph of RMSE with Respect to {}_
    ↪Iterations\n'.format(i),
    ↪'Number of Iterations','Number of_
    ↪Iterations',' $|f(x_{i-1})-f(x_i)|$ ','RMSE')
```


Wine Quality Dataset | Graph of $|f(x_{i-1})-f(x_i)|$ with Respect to 100 Iterations Wine Quality Dataset | Graph of RMSE with Respect to 100 Iterations




Wine Quality Dataset | Graph of $|f(x_{i-1})-f(x_i)|$ with Respect to 500 Iterations Wine Quality Dataset | Graph of RMSE with Respect to 500 Iterations



Wine Quality Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 1000 Iterations



Wine Quality Dataset | Graph of RMSE with Respect to 1000 Iterations




Applying Different Iterations on Parkison Dataset with $\mu=0.00000001$


```
[25]: for i in range(1,imax_d):
        #Calculating new beta using Gradient Descent using  $\mu = 0.00000001$ 
        beta = _
        minimize_GD(parkison_Xtrain,parkison_Ytrain,parkison_Xtest,parkison_Ytest,i,0.
        00000001)

        #Plotting the Loss and RMSE with different Iterations
        plot_loss([j for j in range(i)],loss_difference_values,rmse_values,
        'Parkison Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with Respect to {}_
        Iterations\n'.format(i),
        'Parkison Dataset | Graph of RMSE with Respect to {}_
        Iterations\n'.format(i),
        'Number of Iterations','Number of _
        Iterations',' $|f(x_{i-1}) - f(x_i)|$ ','RMSE')
```


Parkison Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 100 Iterations




Parkison Dataset | Graph of RMSE with Respect to 100 Iterations




Parkison Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 500 Iterations




Parkison Dataset | Graph of RMSE with Respect to 500 Iterations



Parkison Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 1000 Iterations



Parkison Dataset | Graph of RMSE with Respect to 1000 Iterations





Applying Different Iterations on Airfare Demand Dataset with $\mu=0.000000000001$



```
[26]: for i in imax_d:
        #Calculating new beta using Gradient Descent using  $\mu = 0.000000000001$ 
        beta =  $\square$ 
         $\hookrightarrow$ minimize_GD(airfare_Xtrain,airfare_Ytrain,airfare_Xtest,airfare_Ytest,i,0.
         $\hookrightarrow$ 0000000000001)

        #Plotting the Loss and RMSE with different Iterations
        plot_loss([j for j in range(i)],loss_difference_values,rmse_values,
                    'Airfare Demand Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with Respect  $\square$ 
         $\hookrightarrow$ to {} Iterations\n'.format(i),
                    'Airfare Demand Dataset | Graph of RMSE with Respect to {}  $\square$ 
         $\hookrightarrow$ Iterations\n'.format(i),
                    'Number of Iterations','Number of  $\square$ 
         $\hookrightarrow$ Iterations',' $|f(x_{i-1}) - f(x_i)|$ ','RMSE')
```



Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 100 Iterations



Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 500 Iterations



Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to 1000 Iterations

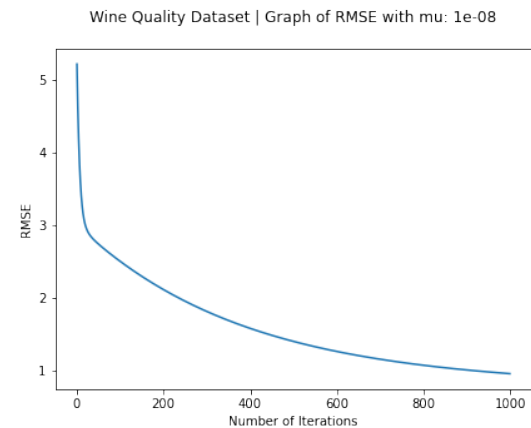

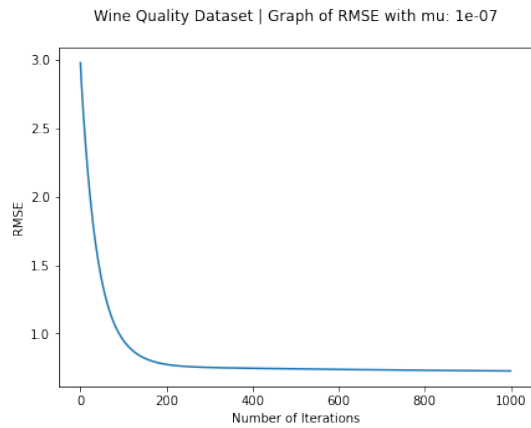
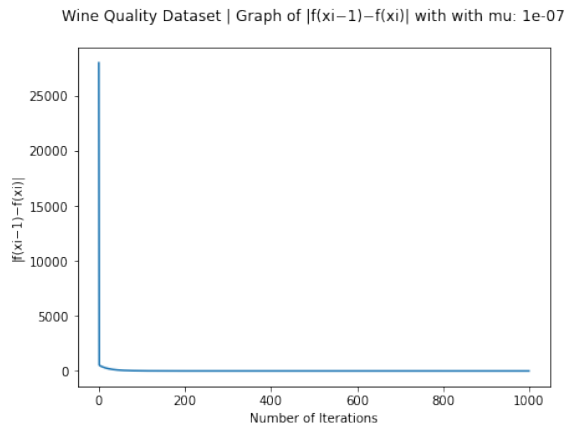


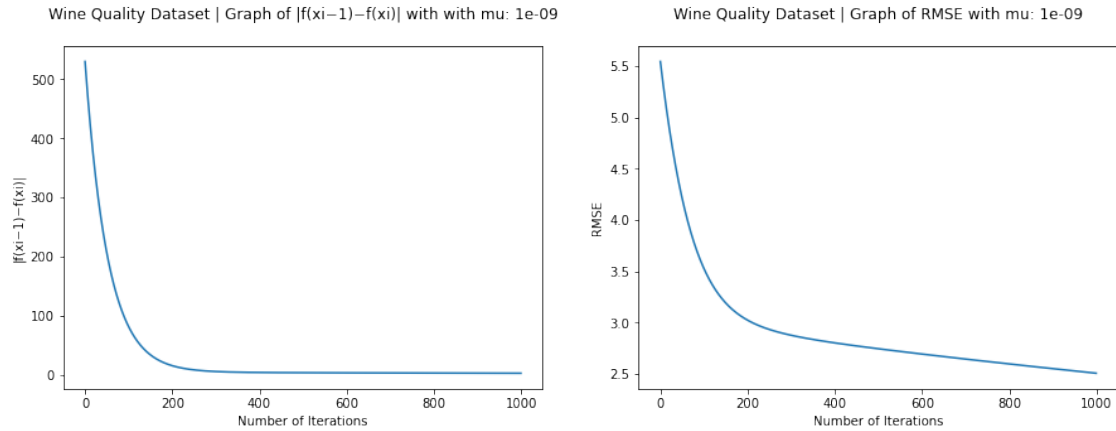
From the above plots and visualization we can see that the minimum loss occurs after 1000 iterations. So selecting $\text{imax} = 1000$ for further experimentation

Applying Different μ values on Wine Quality Dataset with $\text{imax} = 1000$

```
[27]: for mu in mu_wine:
    #Calculating new beta using Gradient Descent using imax = 1000
    beta = minimize_GD(wine_Xtrain,wine_Ytrain,wine_Xtest,wine_Ytest,1000,mu)

    #Plotting the Loss and RMSE with different mu values
    plot_loss([i for i in range(1000)],loss_difference_values,rmse_values,
        'Wine Quality Dataset | Graph of  $|f(x_{i-1})-f(x_i)|$  with with mu: ',
        '\n'.format(mu),
        'Wine Quality Dataset | Graph of RMSE with mu: ', '\n'.format(mu),
        'Number of Iterations', 'Number of ',
        'Iterations', ' $|f(x_{i-1})-f(x_i)|$ ', 'RMSE')
```




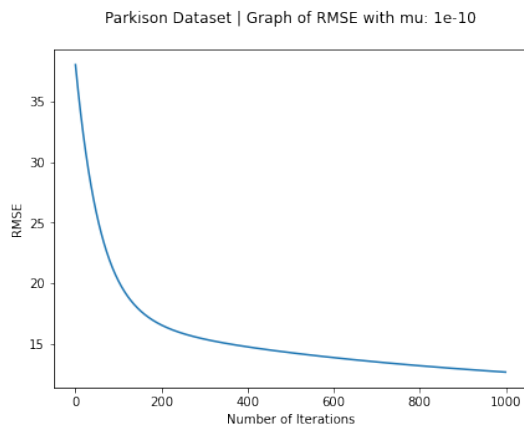
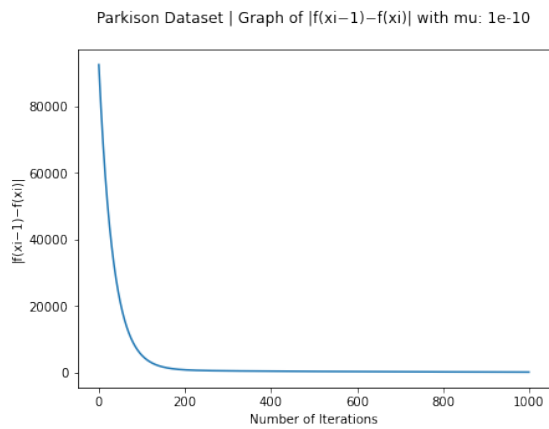
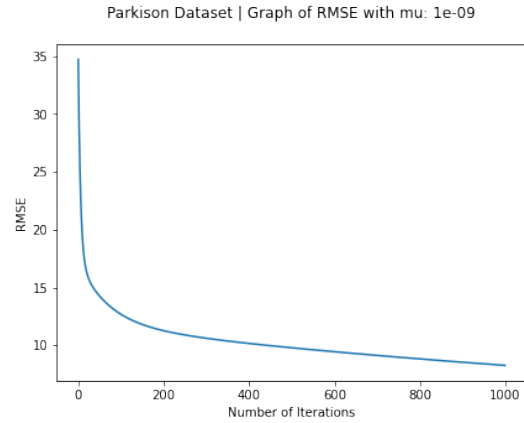
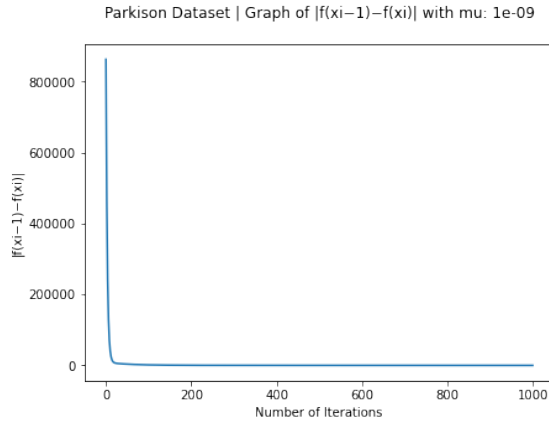


Applying Different μ values on Parkison Dataset with $\text{imax} = 1000$

```
[28]: for mu in mu_parkison:
        #Calculating new beta using Gradient Descent using imax = 1000
        beta =  $\square$ 
         $\rightarrow$ minimize_GD(parkison_Xtrain,parkison_Ytrain,parkison_Xtest,parkison_Ytest,1000,mu)

        #Plotting the Loss and RMSE with different mu values
        plot_loss([i for i in range(1000)],loss_difference_values,rmse_values,
                    'Parkison Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with  $\mu: \{\}$ \n'.
                     $\rightarrow$ format(mu),
                    'Parkison Dataset | Graph of RMSE with  $\mu: \{\}$ \n'.format(mu),
                    'Number of Iterations','Number of  $\square$ 
                     $\rightarrow$ Iterations',' $|f(x_{i-1}) - f(x_i)|$ ','RMSE')
```






Applying Different μ values on Airfare Demand Dataset with $\text{imax} = 1000$


```
[29]: for mu in mu_airfare:
        #Calculating new beta using Gradient Descent using imax = 1000
        beta =  $\square$ 
         $\hookrightarrow$ minimize_GD(airfare_Xtrain,airfare_Ytrain,airfare_Xtest,airfare_Ytest,1000,mu)

        #Plotting the Loss and RMSE with different mu values
        plot_loss([i for i in range(1000)],loss_difference_values,rmse_values,
                    'Airfare Demand Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with  $\mu: \{\}\backslash n'$ .
         $\hookrightarrow$ format(mu),
                    'Airfare Demand Dataset | Graph of RMSE with  $\mu: \{\}\backslash n'$ .format(mu),
                    'Number of Iterations','Number of  $\square$ 
         $\hookrightarrow$ Iterations',' $|f(x_{i-1}) - f(x_i)|$ ','RMSE')
```



Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with $\mu: 1e-12$




Airfare Demand Dataset | Graph of RMSE with $\mu: 1e-12$




Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with $\mu: 1e-13$




Airfare Demand Dataset | Graph of RMSE with $\mu: 1e-13$



Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with $\mu: 1e-14$



Airfare Demand Dataset | Graph of RMSE with $\mu: 1e-14$



0.2 Exercise 3: Steplength Control for Gradient Descent

Function to Plot Loss Difference and RMSE

```
[30]: def plot_stepsize_iteration(xvalues1, xvalues2, loss_difference, rmse,
    →, title_graph1, title_graph2,
    xlabel_graph1, xlabel_graph2,
    →, ylabel_graph1, ylabel_graph2):
    #Plotting Total Iteration vs Loss Difference Graph
    fig = plt.figure(figsize=(15,5))
    ax = fig.add_subplot(121)
    ax.plot(xvalues1, loss_difference)
    ax.set_title(title_graph1)
    ax.set_xlabel(xlabel_graph1)
    ax.set_ylabel(ylabel_graph1)

    #Plotting different MU values vs RMSE Loss
    ax1 = fig.add_subplot(122)
    ax1.plot(xvalues2, rmse)
    ax1.set_title(title_graph2)
    ax1.set_xlabel(xlabel_graph2)
    ax1.set_ylabel(ylabel_graph2)
```

Function to calculate optimum learning rate using Backtracking Algorithm

```
[31]: def stepsize_backtracking(X, Y, B):
    mu = 1          #Starting Learning rate value
    alpha = 0.1     #Value for Alpha
    beta = 0.5      #Value for Beta

    #Iterate until the following condition is meet:
    # $f(x) - f(x - \mu f'(x)) < \mu f'(x)^T f(x)$ 
    while((loss_function(X, Y, B) - loss_function(X, Y, B + (mu * -1 * dL(X, Y, B)))) <
    → (alpha * mu * (dL(X, Y, B).T @ dL(X, Y, B)))):
        mu = mu * beta
    return mu
```

Function to Minimize Gradient Descent using Backtracking algorithm

```
[32]: def minimize_GD_backtracking(X, Y, X_test, Y_test, imax):
    #Creating and Initializing arrays to store loss difference and RMSE
    global loss_difference_values, rmse_values
    loss_difference_values, rmse_values = np.array([]), np.array([])

    #Initializing beta with 0s
    beta = np.zeros(shape=(len(X[0]), 1))
    for i in range(imax):
        #Calculating new Beta using Gradient Descent and backtracking algorithm
        beta_ = beta - stepsize_backtracking(X, Y, beta) * dL(X, Y, beta)
```

```

        #Calculating Loss difference between old Beta and New Beta
        loss_difference_values = np.
        ↪append(loss_difference_values,loss_difference(X,Y,beta,beta_))

        #Calculating RMSE loss based on new Beta
        rmse_values = np.append(rmse_values,rmse(X_test,Y_test,beta_))

        #Copying New beta to Old Beta for Next Iteration
        beta = np.copy(beta_)
        return beta_


```

```

[33]: #Calculating new beta using Gradient Descent using imax = 1000
beta =_
        ↪minimize_GD_backtracking(wine_Xtrain,wine_Ytrain,wine_Xtest,wine_Ytest,1000)

#Plotting the Loss and RMSE with different mu values
plot_stepsize_iteration([i for i in range(1000)],[j for j in_
        ↪range(1000)],loss_difference_values,rmse_values,
        'Wine Quality Dataset | Graph of |f(xi-1)-f(xi)| with 1000_
        ↪Iterations\n'.format(i),
        'Wine Quality Dataset | Graph of RMSE with 1000 Iterations\n'.
        ↪format(i),
        'Number of Iterations','Number of_
        ↪Iterations','|f(xi-1)-f(xi)|','RMSE')

```



```

[34]: #Calculating new beta using Gradient Descent using imax = 1000
beta =_
        ↪minimize_GD_backtracking(parkison_Xtrain,parkison_Ytrain,parkison_Xtest,parkison_Ytest,1000)



#Plotting the Loss and RMSE with different mu values

```

```

plot_stepsize_iteration([i for i in range(1000)], [j for j in
    ↪range(1000)], loss_difference_values, rmse_values,
    'Parkison Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with 1000 Iterations\n'.
    ↪format(i),
    'Parkison Dataset | Graph of RMSE with 1000 Iterations\n'.format(i),
    'Number of Iterations', 'Number of
    ↪Iterations', ' $|f(x_{i-1}) - f(x_i)|$ ', 'RMSE')

```




```

[35]: #Calculating new beta using Gradient Descent using imax = 1000
beta =
    ↪minimize_GD_backtracking(airfare_Xtrain,airfare_Ytrain,airfare_Xtest,airfare_Ytest,1000)


#Plotting the Loss and RMSE with different mu values
plot_stepsize_iteration([i for i in range(1000)], [j for j in
    ↪range(1000)], loss_difference_values, rmse_values,
    'Airfare Demand Dataset | Graph of  $|f(x_{i-1}) - f(x_i)|$  with Respect to
    ↪different Iterations\n'.format(i),
    'Airfare Demand Dataset | Graph of RMSE with Respect to different
    ↪Iterations\n'.format(i),
    'Different Iterations', 'Different
    ↪Iterations', ' $|f(x_{i-1}) - f(x_i)|$ ', 'RMSE')

```

Airfare Demand Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with Respect to different Iterations



Airfare Demand Dataset | Graph of RMSE with Respect to different Iterations



Function to calculate optimum learning rate using Bold Driver Algorithm

```
[36]: def steplength_bolddriver(X,Y,B,alpha_old,alpha_plus,alpha_minus):
    #Increasing alpha value using alpha+
    alpha = alpha_old*alpha_plus

    #Iterating until following condition is meet:
    #f(x) - f(x + μd)  0
    while(loss_function(X,Y,B) - loss_function(X,Y,B - (alpha * dL(X,Y,B))) <= 0):
        #Slowly Increasing alpha using alpha-
        alpha = alpha * alpha_minus
    return alpha
```

Function to Minimize Gradient Descent using Bold Driver algorithm

```
[37]: def minimize_GD_bolddriver(X,Y,X_test,Y_test,imax):
    #Creating and Initializing arrays to store loss difference and RMSE
    global loss_difference_values , rmse_values
    loss_difference_values , rmse_values = np.array([]) , np.array([])

    #Initializing beta with 0s
    theta = np.zeros(shape=(len(X[0]),1))
    alpha = 1
    for i in range(imax):
        #Calculating alpha based bold driver algorithm
        alpha = steplength_bolddriver(X,Y,theta,alpha,1.1,0.5)

        #Calculating new Beta using Gradient Descent and backtracking algorithm
        theta_ = theta - alpha * dL(X,Y,theta)

        #Calculating Loss difference between old Beta and New Beta
        loss_difference_values = np.
        append(loss_difference_values,loss_difference(X,Y,theta,theta_))
```

```

#Calculating RMSE loss based on new Beta
rmse_values = np.append(rmse_values,rmse(X_test,Y_test,theta_))

#Copying New beta to Old Beta for Next Iteration
theta = np.copy(theta_)
return theta_

```


```

[38]: #Calculating new beta using Gradient Descent using imax = 1000
beta =_
↳minimize_GD_bolddriver(wine_Xtrain,wine_Ytrain,wine_Xtest,wine_Ytest,1000)


#Plotting the Loss and RMSE with different mu values
plot_stepsize_iteration([i for i in range(1000)],[j for j in_
↳range(1000)],loss_difference_values,rmse_values,
    'Wine Quality Dataset | Graph of |f(xi-1)-f(xi)| with 1000_
↳Iterations\n'.format(i),
    'Wine Quality Dataset | Graph of RMSE with 1000 Iterations\n'.
↳format(i),
    'Number of Iterations','Number of_
↳Iterations','|f(xi-1)-f(xi)|','RMSE')

```

Wine Quality Dataset | Graph of $|f(x_{i-1}) - f(x_i)|$ with 1000 Iterations



Wine Quality Dataset | Graph of RMSE with 1000 Iterations



```

[39]: #Calculating new beta using Gradient Descent using imax = 1000
beta =_
↳minimize_GD_bolddriver(parkison_Xtrain,parkison_Ytrain,parkison_Xtest,parkison_Ytest,1000)



#Plotting the Loss and RMSE with different mu values
plot_stepsize_iteration([i for i in range(1000)],[j for j in_
↳range(1000)],loss_difference_values,rmse_values,
    'Parkison Dataset | Graph of |f(xi-1)-f(xi)| with 1000 Iterations\n'.
↳format(i),

```

```

'Parkison Dataset | Graph of RMSE with 1000 Iterations\n'.format(i),
'Number of Iterations','Number of_
↪Iterations','|f(xi-1)-f(xi)|','RMSE')

```





```

[40]: #Calculating new beta using Gradient Descent using imax = 1000
beta =_
↪minimize_GD_bolddriver(airfare_Xtrain,airfare_Ytrain,airfare_Xtest,airfare_Ytest,1000)

#Plotting the Loss and RMSE with different mu values
plot_stepsize_iteration([i for i in range(1000)], [j for j in_
↪range(1000)], loss_difference_values, rmse_values,
'Airfare Demand Dataset | Graph of |f(xi-1)-f(xi)| with 1000_
↪Iterations\n'.format(i),
'Airfare Demand Dataset | Graph of RMSE with 1000 Iterations\n'.
↪format(i),
'Number of Iterations','Number of_
↪Iterations','|f(xi-1)-f(xi)|','RMSE')

```



Based on Backtracking and Bold Driver Algorithm performance, I think Bold driver did a bit better job as compared to Backtracking as it converges faster than the later But with a very smaller margin. The margin is so small that we can use any of these algorithm without any significant performance drop.