

# Lab 10

January 29, 2022

Machine Learning Lab

Lab 10

## 0.0.1 Importing Packages

```
[1]: import pandas as pd                #Importing Pandas
import numpy as np                    #Importing Numpy
import matplotlib.pyplot as plt      #Importing Matplotlib
from sklearn.decomposition import NMF #Importing SKlearn Non-
    ↪ Negative Matrix Factorization
from sklearn.model_selection import GridSearchCV #Importing GridSearchCV
import warnings                      #Importing Warnings
warnings.filterwarnings('ignore')
```

## 0.0.2 Exercise 1: Exploring Movie Recommendation Dataset

Reading Datasets

Reading the dataset file containing Users and Items references

```
[2]: u_data = pd.read_csv('ml-100k/ml-100k/u.data', sep='\t', names=['user id',
    ↪ 'item id', 'rating', 'timestamp'], header=None)
u_data.head()
```

```
[2]:   user id  item id  rating  timestamp
0      196     242        3   881250949
1      186     302        3   891717742
2       22     377        1   878887116
3      244       51        2   880606923
4      166     346        1   886397596
```

Reading the Users Data

```
[3]: u_user = pd.read_csv('ml-100k/ml-100k/u.user', sep='|', names=['user id',
    ↪ 'age', 'gender', 'occupation', 'zip code'])
u_user.head()
```

```
[3]:   user id  age  gender  occupation  zip code
0         1   24      M  technician    85711
```

1	2	53	F	other	94043
2	3	23	M	writer	32067
3	4	24	M	technician	43537
4	5	33	F	other	15213

## Reading the Genre Data

```
[4]: u_genre = pd.read_csv('ml-100k/ml-100k/u.genre', sep='|', names=['genre', 'id'])
u_genre.head()
```

```
[4]:      genre  id
0    unknown  0
1     Action  1
2  Adventure  2
3  Animation  3
4 Children's  4
```

## Reading the Items/Movies Data

```
[5]: u_item = pd.read_csv('ml-100k/ml-100k/u.item', encoding='latin-1', sep='|',
    ↪names=['movie id', 'movie title', 'release date', 'video release date', 'IMDb_
    ↪URL', 'unknown', 'Action', 'Adventure', 'Animation', 'Children\'s', 'Comedy', 'Crime', 'Documentary
u_item.head()
```

```
[5]:  movie id      movie title release date  video release date  \
0      1  Toy Story (1995)  01-Jan-1995      NaN
1      2  GoldenEye (1995)  01-Jan-1995      NaN
2      3  Four Rooms (1995)  01-Jan-1995      NaN
3      4  Get Shorty (1995)  01-Jan-1995      NaN
4      5  Copycat (1995)  01-Jan-1995      NaN

      IMDb URL  unknown  Action  \
0  http://us.imdb.com/M/title-exact?Toy%20Story%2...      0      0
1  http://us.imdb.com/M/title-exact?GoldenEye%20( ...      0      1
2  http://us.imdb.com/M/title-exact?Four%20Rooms%...      0      0
3  http://us.imdb.com/M/title-exact?Get%20Shorty%...      0      1
4  http://us.imdb.com/M/title-exact?Copycat%20(1995)      0      0

      Adventure  Animation  Children's  ...  Fantasy  Film-Noir  Horror  Musical  \
0      0      1      1  ...      0      0      0      0
1      1      0      0  ...      0      0      0      0
2      0      0      0  ...      0      0      0      0
3      0      0      0  ...      0      0      0      0
4      0      0      0  ...      0      0      0      0

      Mystery  Romance  Sci-Fi  Thriller  War  Western
0      0      0      0      0      0      0
1      0      0      0      1      0      0
2      0      0      0      1      0      0
```


3	0	0	0	0	0	0
4	0	0	0	1	0	0

[5 rows x 24 columns]

### Showcasing how the ratings vary across users

```
[6]: #For finding total ratings across user,
#First we group the data based on user Id and then count the total rows which
    ↳will indicate total ratings given by users
rating_across_users = u_data.groupby(by = 'user id')['rating'].count().
    ↳sort_values(ascending=True)

#Plotting the Graph: Users vs Number of Ratings
fig = plt.figure(figsize=(18,8))
plt.bar(rating_across_users.index, rating_across_users.values)
plt.title('How the ratings vary across users')
plt.xlabel('Users')
plt.ylabel('Number of Ratings')
plt.show()
```

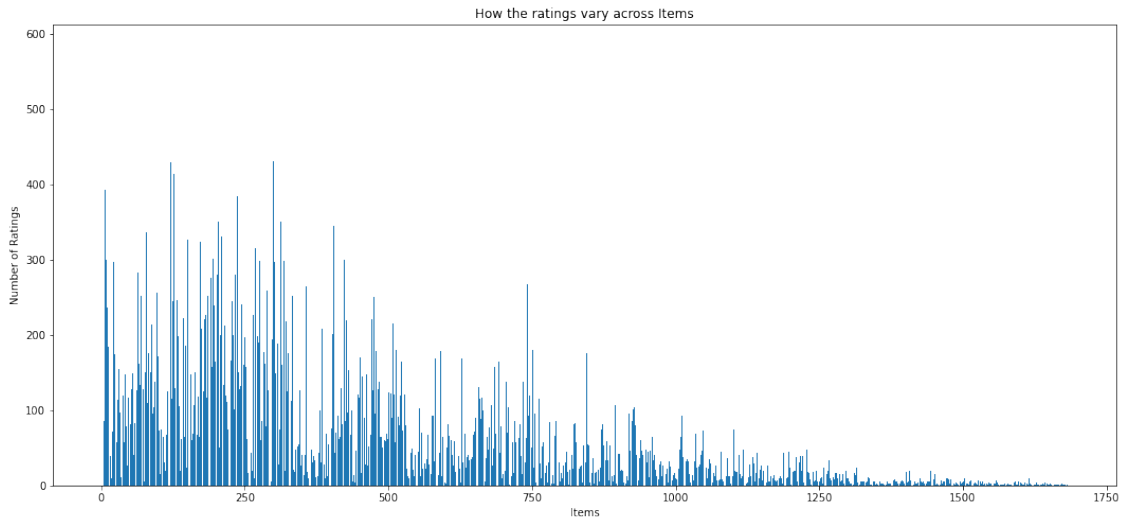


### Showcase how the ratings vary across items

```
[7]: #For finding total ratings across Items,
#First we group the data based on item Id and then count the total rows which
    ↳will indicate total ratings given for Items
rating_across_items = u_data.groupby(by='item id')['rating'].count().
    ↳sort_values(ascending = True)

#Plotting the Graph: Users vs Number of Ratings
```

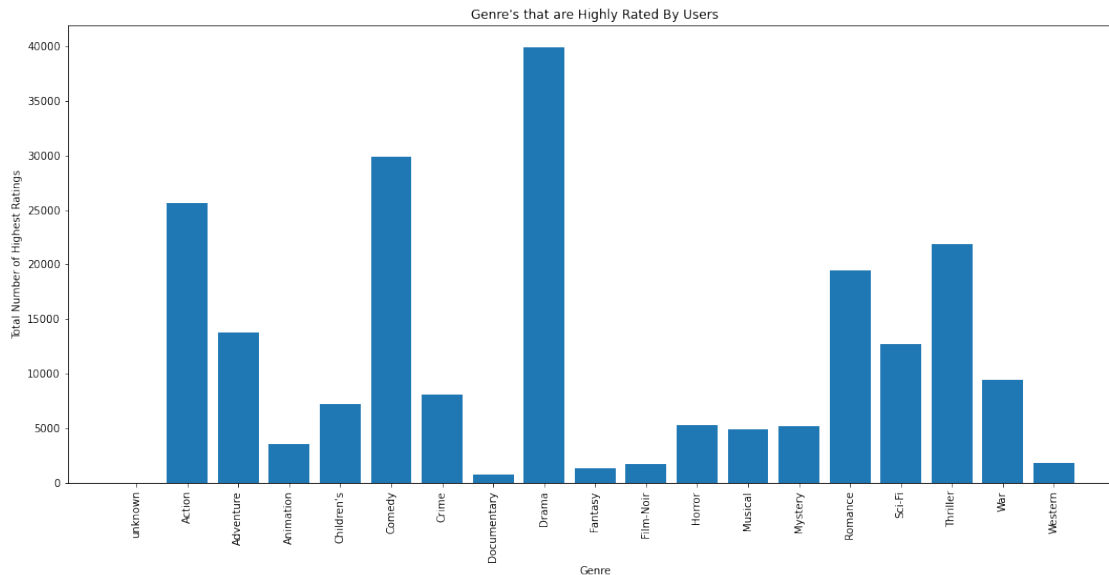
```
fig = plt.figure(figsize=(18,8))
plt.bar(rating_across_items.index, rating_across_items.values)
plt.title('How the ratings vary across Items')
plt.xlabel('Items')
plt.ylabel('Number of Ratings')
plt.show()
```



Are there genres that are more highly rated than others?

```
[8]: #First Extracting only that rows that has the maximum ratings i.e. 5
#Then to find the genre and there count, we can simply sum up all the rows
    ↳ horizontally and the result will show there count individually
highest_rated_items = u_data[u_data['rating'] == max(u_data['rating'])]
highest_rated_items = pd.merge(u_data, u_item, left_on='item id',
    ↳ right_on='movie id')
highest_rated_items = (highest_rated_items[list(u_genre['genre'])] == 1).
    ↳ sum(axis=0)

#Plotting the graph indicating total ratings for each genre
fig = plt.figure(figsize=(18,8))
plt.bar(highest_rated_items.index, highest_rated_items.values)
plt.title('Genre\'s that are Highly Rated By Users')
plt.xlabel('Genre')
plt.ylabel('Total Number of Highest Ratings')
plt.xticks(rotation = 90)
plt.show()
```



What age groups prefer what genres based on ratings?

```
[9]: #Defining different age brackets with a difference of 10 each
age_bracket = np.arange(0,80,10)
```

```
[10]: #Creating a dictionary to hold best genre for each age group
aged_group_preference = {}

#Extracting rows in which the item highest Ratings i.e. 5
merged_df = u_data[u_data['rating'] == max(u_data['rating'])]

#Merging the datasets
merged_df = pd.merge(merged_df, u_item, left_on='item id', right_on='movie id')
merged_df = pd.merge(merged_df, u_user, on='user id')

#Iterating over all the age groups
for age in age_bracket:
    #Extracting ratings which are given by users in the specified age bracket
    aged_df = merged_df[merged_df['age'].between(age, age+10)]

    #Summing the Genre count and getting the genre with the maximum count and
    ↪ saving it into the dictionary
    aged_group_preference['{}-{}'.format(age, age+10)] =
    ↪ (aged_df[list(u_genre['genre'])] == 1).sum(axis=0).idxmax(axis = 0)

#Converting the Dictionary into a Dataframe for Displaying the results
aged_df = pd.DataFrame.from_dict(aged_group_preference, orient='index').
    ↪ reset_index()
```

```
aged_df.columns = ['Age Group', 'Popular Genre']
aged_df
```

```
[10]:  Age Group Popular Genre
      0      0-10      Comedy
      1     10-20      Drama
      2     20-30      Drama
      3     30-40      Drama
      4     40-50      Drama
      5     50-60      Drama
      6     60-70      Drama
      7     70-80      Drama
```

### 0.0.3 Exercise 2: Implementing basic matrix factorization (MF) technique for recommender systems

```
[11]: #Initializing the random seed value
      random_seed = 3116
```

#### Normalizing the data

```
[12]: def normalize_matrix(matrix):
      #Normalizing the dataset by:  $(X - X.mean)/X.standard\_Deviation$ 
      return (matrix - matrix.mean())/matrix.std()
```

#### Function to transform the dataset into a Rating Matrix of Users and Item

```
[13]: def extract_rating_matrix(dataset):
      #Initializing the rating matrix with zeros
      rating_matrix = np.zeros(shape=(number_of_users, number_of_items))

      #Iterating over all the rows of the dataset
      for row in dataset:
          #The row represent the users and the column represent the items and the
          ↳corresponding value represent the rating for that item
          rating_matrix[row[0]-1, row[1]-1] = row[2]

      #Returning the Rating matrix
      return rating_matrix
```

#### Function to perform K-fold cross validation and return the training and the validation sets

```
[14]: def k_fold_cross_validation(dataset, k, K):
      #Checking if the given k is larger than the allowed K partitions
      if k >= K:
          raise Exception('The Requested k-Fold should not be Greater than total
          ↳Folds')
```

```

    #Calculating the Validation set size based on dataset size and total
    ↪partitions given
    set_size = int(len(dataset)/K)

    #Splitting the Dataset into Training and Validation sets based on set size.
    ↪The size
    #set size calculated above and all the remaining data points after removing
    ↪test
    dataset_train = np.delete(dataset,[i for i in
    ↪range(k*set_size,(k*set_size)+set_size)],axis=0)
    dataset_validation = dataset[k*set_size:(k*set_size)+set_size,:]

    #Returning the created Training and Validation datasets.
    return dataset_train, dataset_validation

```

Function to calculate the RMSE between the actual rating matrix and the predicted rating matrix

```

[15]: def rmse(rating_matrix, p , q):
    #Returning the RMSE
    return np.sqrt(np.divide(np.sum(np.square(np.subtract(rating_matrix, np.
    ↪dot(p, q))))),rating_matrix.shape[0] * rating_matrix.shape[1]))

```

Function which implements Stochastic Gradient Descent (SGD) for learning the matrix P and Q

```

[16]: def learn_latent_factors(dataset_train, dataset_test, lamda, alpha, K):
    #Extracting and normalizing the Rating matrix of the Training dataset
    r = normalize_matrix(extract_rating_matrix(dataset_train))

    #Extracting and normalizing the Rating matrix of the Test dataset
    r_v = normalize_matrix(extract_rating_matrix(dataset_test))

    #Intializing the random P and Q matrix for matrix factorization, there
    ↪dimensions are as follows:
    #P -> (# of users, K latent Features)
    #Q -> (K latent Features, # of Items)
    p = np.random.normal(normal_dis_mean, normal_dis_std, size =
    ↪(number_of_users, K))
    q = np.random.normal(normal_dis_mean, normal_dis_std, size = (K,
    ↪number_of_items))

    #Intializing lists for storing the train and test RMSE in each Epochs
    train_rmse = np.array([])
    validation_rmse = np.array([])

    #Learning the matrix P and Q for specified number of Epochs
    for i in range(epochs):

```

```

#Iterating over all the rows in the training dataset
for row in dataset_train:
    #Extracting the User Id and Item Id from the row
    user_id, item_id = row[0] - 1, row[1] - 1

    #Calculating the difference between the actual rating and the
    ↪ predicted rating
    e = r[user_id, item_id] - np.dot(p[user_id,:], q[:,item_id])

    #Iterating over all the latent feature
    for k in range(K):
        #Updating the values of  $P_{u,k}$  and  $Q_{i,k}$ 
        p[user_id, k] = p[user_id, k] + alpha * (e * q[k, item_id] -
    ↪ lamda * p[user_id, k])
        q[k, item_id] = q[k, item_id] + alpha * (e * p[user_id, k] -
    ↪ lamda * q[k, item_id])

    #Appending the Train and Validating RMSE in the list
    train_rmse = np.append(train_rmse, rmse(r, p, q))
    validation_rmse = np.append(validation_rmse, rmse(r_v, p, q))

    #Checking the Stopping Condition, Stopping the Learning process when the
    ↪ difference between the last two epoch is negligible
    if i > 1 and (validation_rmse[-2] - validation_rmse[-1] < 0.001):
        break

    #Returning the calculated P, Q, Train RMSE and Validation RMSE
    return p, q, train_rmse, validation_rmse

```

Optimizing the hyper-parameters i.e. regularization constant, learning rate, k latent dimensions

Initializing variables to be used by different functions

```

[17]: #Number of Unique Users in the Dataset
number_of_users = len(u_user)

#Number of Unique Items in the Dataset
number_of_items = len(u_item)

#Initializing the Training set Size
train_ratio = 0.9

```

Initializing variables storing the values of parameters of different algorithm

```

[18]: #Number of Epochs
epochs = 10

```



```

#Mean for Normal Distribution samples
normal_dis_mean = 0

#Standard Deviation for Normal Distribution samples
normal_dis_std = 1

#Number of Folds for K fold cross validation
K_folds = 5

```

### Initializing different Hyperparameters sets for Hyperparameter Optimization

```

[19]: lamdas = [0.1, 0.2, 0.3]
      alphas = [0.005, 0.003, 0.001]
      latent_k = [10,20,30]

```

### Merging and Splitting the Dataset

```

[20]: #Merging the u.data and u.item dataset
      merged_df = pd.merge(u_data, u_item, left_on='item id', right_on='movie id')

      #Merging the merged dataset with u.users
      merged_df = pd.merge(merged_df, u_user, on='user id')

```

```

[21]: #Extracting the rows randomly from the merged dataset for Training set
      merged_train = merged_df.sample(frac=train_ratio)

      #Dropping that rows which are included in training dataset
      merged_test = merged_df.drop(merged_train.index)

      #Converting the dataset into numpy arrays
      merged_train = merged_train.to_numpy()
      merged_test = merged_test.to_numpy()

```

### Performing the Hyperparameter Optimization

```

[22]: #Initializing a matrix for storing the training RMSE for each Hyperparameter_
      ↪ combination
      hyperparameters_train_rmse = np.
      ↪ zeros(shape=(len(lamdas),len(alphas),len(latent_k)))

      #Initializing a matrix for storing the Validating RMSE for each Hyperparameter_
      ↪ combination
      hyperparameters_validation_rmse = np.
      ↪ zeros(shape=(len(lamdas),len(alphas),len(latent_k)))

```

```

[23]: #Iterating for all values of Lambdas
      for index_l, lamda in enumerate(lamdas):

          #Iterating for all values of Alphas

```

```

for index_a, alpha in enumerate(alphas):

    #Iterating for all values of Latent Features
    for index_lk, l_k in enumerate(latent_k):

        #Initializing the average Training and Validating RMSE for each K
        ↪fold
        average_train_rmse, average_validation_rmse = 0 , 0

        #Iterating for all values of K fold
        for k in range(K_folds):

            #Splitting the training dataset into Train and Validation using
            ↪K fold Cross validation
            train, validation = k_fold_cross_validation(merged_train, k,
            ↪K_folds)

            #Applying the SGD for Calculating the matrix P and Q using a
            ↪hyperparameter combination
            p, q, train_rmse, validation_rmse =
            ↪learn_latent_factors(merged_train, validation, lamda, alpha, l_k)

            #Adding the Training and Validation RMSE
            average_train_rmse += train_rmse[-1]
            average_validation_rmse += validation_rmse[-1]

            #Averaging the values for Train and Validation RMSE
            average_train_rmse, average_validation_rmse = average_train_rmse/
            ↪K_folds, average_validation_rmse/K_folds

            #Printing the Output of RMSE for selected Hyperparameter Combination
            print('Hyperparameter Selected: ')
            print('Lambda: {} \t Alpha: {} \t Latent Features: {}'.format(lamda,
            ↪alpha, l_k))
            print('Training RMSE: {} \t Validation RMSE: {}'.format(
            ↪format(average_train_rmse, average_validation_rmse))

            #Saving the calculated RMSE in the list of both Training and
            ↪Validation
            hyperparameters_train_rmse[index_l,index_a,index_lk] =
            ↪average_train_rmse
            hyperparameters_validation_rmse[index_l,index_a,index_lk] =
            ↪average_validation_rmse

```

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.005      Latent Features: 10

Training RMSE: 2.868509893658427      Validation RMSE: 2.9597379940278317

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.005      Latent Features: 20

Training RMSE: 3.342494949756881      Validation RMSE: 3.426598602951947

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.005      Latent Features: 30

Training RMSE: 3.6669460422303857      Validation RMSE: 3.7275822539444823

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.003      Latent Features: 10

Training RMSE: 2.6233389226393014      Validation RMSE: 2.72348065261489

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.003      Latent Features: 20

Training RMSE: 3.2864389817621222      Validation RMSE: 3.3727021002841555

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.003      Latent Features: 30

Training RMSE: 3.750172900671404      Validation RMSE: 3.817480761038103

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 10

Training RMSE: 2.458437646820484      Validation RMSE: 2.51997618557773

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 20

Training RMSE: 3.293304418318679      Validation RMSE: 3.366715418916641

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 30

Training RMSE: 3.9002412563875226      Validation RMSE: 3.977953930799921

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.005      Latent Features: 10

Training RMSE: 2.736373374601422      Validation RMSE: 2.830550185541351

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.005      Latent Features: 20

Training RMSE: 3.20244993547063      Validation RMSE: 3.286403732468172

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.005      Latent Features: 30

Training RMSE: 3.3908482812611154      Validation RMSE: 3.4498445597873904

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.003      Latent Features: 10

Training RMSE: 2.4924098604959584      Validation RMSE: 2.592825962359919  
 Hyperparameter Selected:  
 Lambda: 0.2      Alpha: 0.003      Latent Features: 20  
 Training RMSE: 3.1436635614832955      Validation RMSE: 3.2291171435295105  
 Hyperparameter Selected:  
 Lambda: 0.2      Alpha: 0.003      Latent Features: 30  
 Training RMSE: 3.474573512313772      Validation RMSE: 3.542423211713325  
 Hyperparameter Selected:  
 Lambda: 0.2      Alpha: 0.001      Latent Features: 10  
 Training RMSE: 2.380917903078778      Validation RMSE: 2.4444837319566686  
 Hyperparameter Selected:  
 Lambda: 0.2      Alpha: 0.001      Latent Features: 20  
 Training RMSE: 3.1564327820813505      Validation RMSE: 3.2308621391027303  
 Hyperparameter Selected:  
 Lambda: 0.2      Alpha: 0.001      Latent Features: 30  
 Training RMSE: 3.6973659274472945      Validation RMSE: 3.775103542152378  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.005      Latent Features: 10  
 Training RMSE: 2.645282438520627      Validation RMSE: 2.7380224819409738  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.005      Latent Features: 20  
 Training RMSE: 3.034819788659462      Validation RMSE: 3.1202311240426512  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.005      Latent Features: 30  
 Training RMSE: 3.1882769363770995      Validation RMSE: 3.247268782831602  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.003      Latent Features: 10  
 Training RMSE: 2.363905106792751      Validation RMSE: 2.462203321638701  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.003      Latent Features: 20  
 Training RMSE: 3.0188006023682847      Validation RMSE: 3.1057964678622327  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.003      Latent Features: 30  
 Training RMSE: 3.2893276953753534      Validation RMSE: 3.355826853021143  
 Hyperparameter Selected:  
 Lambda: 0.3      Alpha: 0.001      Latent Features: 10

Training RMSE: 2.2769188343833044      Validation RMSE: 2.3355294300340343

Hyperparameter Selected:

Lambda: 0.3      Alpha: 0.001      Latent Features: 20

Training RMSE: 3.0163392379154494      Validation RMSE: 3.0984016342988783

Hyperparameter Selected:

Lambda: 0.3      Alpha: 0.001      Latent Features: 30

Training RMSE: 3.5067156353032347      Validation RMSE: 3.5857673503648924

## Computing the Validation RMSE

### Finding the Best Hyperparameter Combination

```
[24]: #Finding the index containing the minimum RMSE
i,j,k = np.where(hyperparameters_validation_rmse ==
    ↳hyperparameters_validation_rmse.min())

#Printing the Results
print('Best Hyperparameter with the Minimum Validation RMSE:')
print('Lambda: {} \t Alpha: {} \t Latent Features: {}'.format(lamdas[i.
    ↳item()], alphas[j.item()], latent_k[k.item()]))
print('Validation RMSE: {}'.format(hyperparameters_validation_rmse[i,j,k].
    ↳item()))
```

Best Hyperparameter with the Minimum Validation RMSE:

Lambda: 0.3      Alpha: 0.001      Latent Features: 10

Validation RMSE: 2.3355294300340343

## Computing the Test RMSE

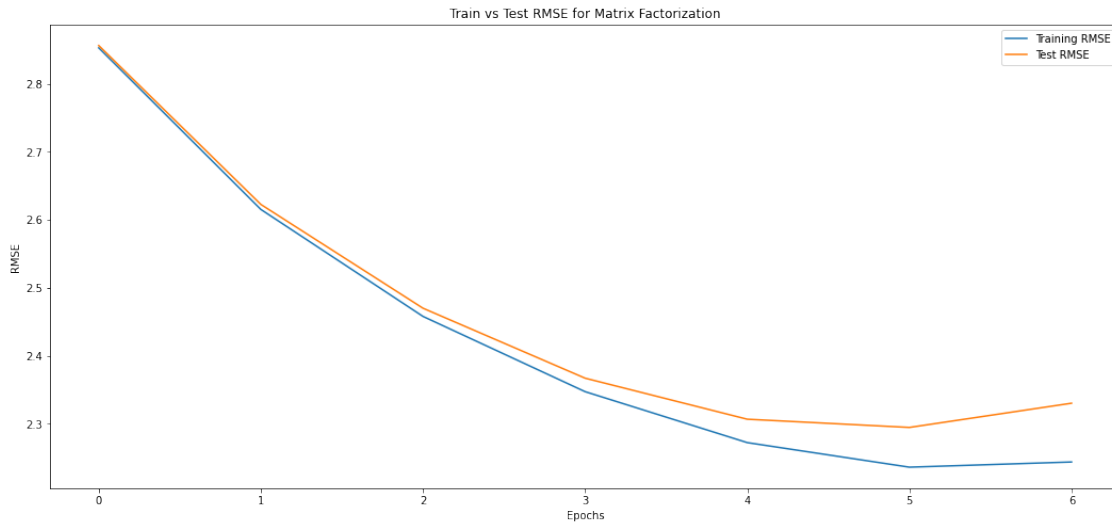
```
[25]: #Performing Matrix Factorization on the Test Dataset with Best Combination of
    ↳Hyperparameters
best_lambda, best_alpha, best_latent_k = lamdas[i.item()], alphas[j.item()],
    ↳latent_k[k.item()]
p, q, train_rmse, test_rmse = learn_latent_factors(merged_train, merged_test,
    ↳best_lambda, best_alpha, best_latent_k)
print('The Test RMSE: {}'.format(test_rmse[-1]))
```

The Test RMSE: 2.3302891774428147

## Plotting the Training RMSE vs Test RMSE for each Epochs

```
[26]: fig = plt.figure(figsize=(18,8))
plt.plot([i for i in range(len(train_rmse))], train_rmse, label='Training RMSE')
plt.plot([i for i in range(len(test_rmse))], test_rmse, label='Test RMSE')
plt.title('Train vs Test RMSE for Matrix Factorization')
plt.xlabel('Epochs')
plt.ylabel('RMSE')
```

```
plt.legend()
plt.show()
```



### 0.0.4 Exercise 3: Recommender Systems using matrix factorization sklearn

#### Performing the Hyperparameter Optimization using Grid Search

```
[27]: #Initializing a matrix for storing the training RMSE for each Hyperparameter
      ↪ combination
hyperparameters_train_rmse = np.zeros(shape=(len(lamdas),len(latent_k)))

#Initializing a matrix for storing the Validating RMSE for each Hyperparameter
      ↪ combination
hyperparameters_validation_rmse = np.zeros(shape=(len(lamdas),len(latent_k)))
```

```
[28]: #Iterating for all values of Lambdas
      for index_l, lamda in enumerate(lamdas):

          #Iterating for all values of Latent Features
          for index_lk, l_k in enumerate(latent_k):

              #Initializing the average Training and Validating RMSE for each K fold
              average_train_rmse, average_validation_rmse = 0 , 0

              #Iterating for all values of K fold
              for k in range(K_folds):

                  #Splitting the training dataset into Train and Validation using K
                  ↪ fold Cross validation
```

```

train, validation = k_fold_cross_validation(merged_train, k,
↳K_folds)

#Using Sklearn NMF function to Matrix Factorization
model = NMF(init='random',n_components = l_k, alpha=lamda,
↳random_state=random_seed)

#Extracting matrix W after fitting the Training Rating Matrix
W = model.fit_transform(extract_rating_matrix(train))

#Extracting matrix H after fitting the Training Rating Matrix
H = model.components_

#Adding the Training and Validation RMSE
average_train_rmse += rmse(extract_rating_matrix(train), W, H)
average_validation_rmse += rmse(extract_rating_matrix(validation),
↳W, H)

#Averaging the values for Train and Validation RMSE
average_train_rmse, average_validation_rmse = average_train_rmse/
↳K_folds, average_validation_rmse/K_folds

#Printing the Output of RMSE for selected Hyperparameter Combination
print('Hyperparameter Selected: ')
print('Lambda: {} \t Alpha: {} \t Latent Features: {}'.format(lamda, alpha,
↳l_k))
print('Training RMSE: {} \t Validation RMSE: {} \n'.
↳format(average_train_rmse, average_validation_rmse))

#Saving the calculated RMSE in the list of both Training and Validation
hyperparameters_train_rmse[index_l,index_lk] = average_train_rmse
hyperparameters_validation_rmse[index_l,index_lk] =
↳average_validation_rmse

```

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 10

Training RMSE: 0.6455496004386772      Validation RMSE: 0.5209734831178708

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 20

Training RMSE: 0.6261437430899222      Validation RMSE: 0.5438638955517104

Hyperparameter Selected:

Lambda: 0.1      Alpha: 0.001      Latent Features: 30

Training RMSE: 0.6109459061417086      Validation RMSE: 0.5647186020802145

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.001      Latent Features: 10  
Training RMSE: 0.6455410880933105      Validation RMSE: 0.5208110910569509

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.001      Latent Features: 20  
Training RMSE: 0.6261731256168357      Validation RMSE: 0.543679951253973

Hyperparameter Selected:

Lambda: 0.2      Alpha: 0.001      Latent Features: 30  
Training RMSE: 0.6110123444640454      Validation RMSE: 0.5642144169255127

Hyperparameter Selected:

Lambda: 0.3      Alpha: 0.001      Latent Features: 10  
Training RMSE: 0.6455446309854551      Validation RMSE: 0.5206179302358835

Hyperparameter Selected:

Lambda: 0.3      Alpha: 0.001      Latent Features: 20  
Training RMSE: 0.6261767517512596      Validation RMSE: 0.5435347118224222

Hyperparameter Selected:

Lambda: 0.3      Alpha: 0.001      Latent Features: 30  
Training RMSE: 0.6109843575875654      Validation RMSE: 0.5640316855817946

## Computing the Validation RMSE

### Finding the Best Hyperparameter Combination

```
[29]: #Finding the index containing the minimum RMSE
i,j = np.where(hyperparameters_validation_rmse ==
    ↳hyperparameters_validation_rmse.min())

#Printing the Results
print('Best Hyperparameter with the Minimum Validation RMSE:')
print('Lambda: {} \t Latent Features: {}'.format(lamdas[i.item()],latent_k[j.
    ↳item()])))
print('Validation RMSE: {}'.format(hyperparameters_validation_rmse[i,j].item()))
```

Best Hyperparameter with the Minimum Validation RMSE:

Lambda: 0.3      Latent Features: 10  
Validation RMSE: 0.5206179302358835

## Computing the Test RMSE

```
[30]: #Performing Matrix Factorization on the Test Dataset with Best Combination of
    ↳Hyperparameters
test_model = NMF(init='random',n_components = 10, alpha=0.3,
    ↳random_state=random_seed)
```



```
#Extracting matrix W after fitting the Training Rating Matrix
W = test_model.fit_transform(extract_rating_matrix(merged_test))

#Extracting matrix H after fitting the Training Rating Matrix
H = test_model.components_
```

```
[31]: print('Test RMSE: {}'.format(rmse(extract_rating_matrix(merged_test), W, H)))
```

```
Test RMSE: 0.27566461477377874
```

From the above results, we can see that the RMSE of the Test set from my implementation of SGD came out to be around 2.33 whereas the RMSE of Test set from the Sklearn implementation of NMF came out to be 0.27, which is a huge improvement. Furthermore, I also observe that the time taken to run every epoch was also very large as compared to the Sklearn implementation which makes it more interesting.